AWS for Data Analytics: Elastic MapReduce

By Jennifer Rubinovitz


@rubinovitz


j@rubinovitz.com

http://bit.ly/hadoopforhackers

Me

- ☐ Rutgers alumnus ‘13
- ☐ hackNY fellow ‘12 + hackNY mentor ‘13

Presently:

- ☐ C.S. M.S. student at Columbia working on Lean Workbench to quantify early stage startups

Outline

- ☐ What is MapReduce
- ☐ What is Hadoop
- ☐ What is Hadoop on AWS
- ☐ Example

Google’s Bluffs Data Center

Motivation: For This Talk

- Can we democratize data science yet?
- EMR is elastic and cheap
- More online + offline resources than ever

Motivation: For MapReduce

- You have data so big you need to parallelize processing (e.g. Google's Index of the World Wide Web)
- Since you have so many nodes you need to assume there will always be failures

<u>Solution:</u> Google MapReduce and Google File System

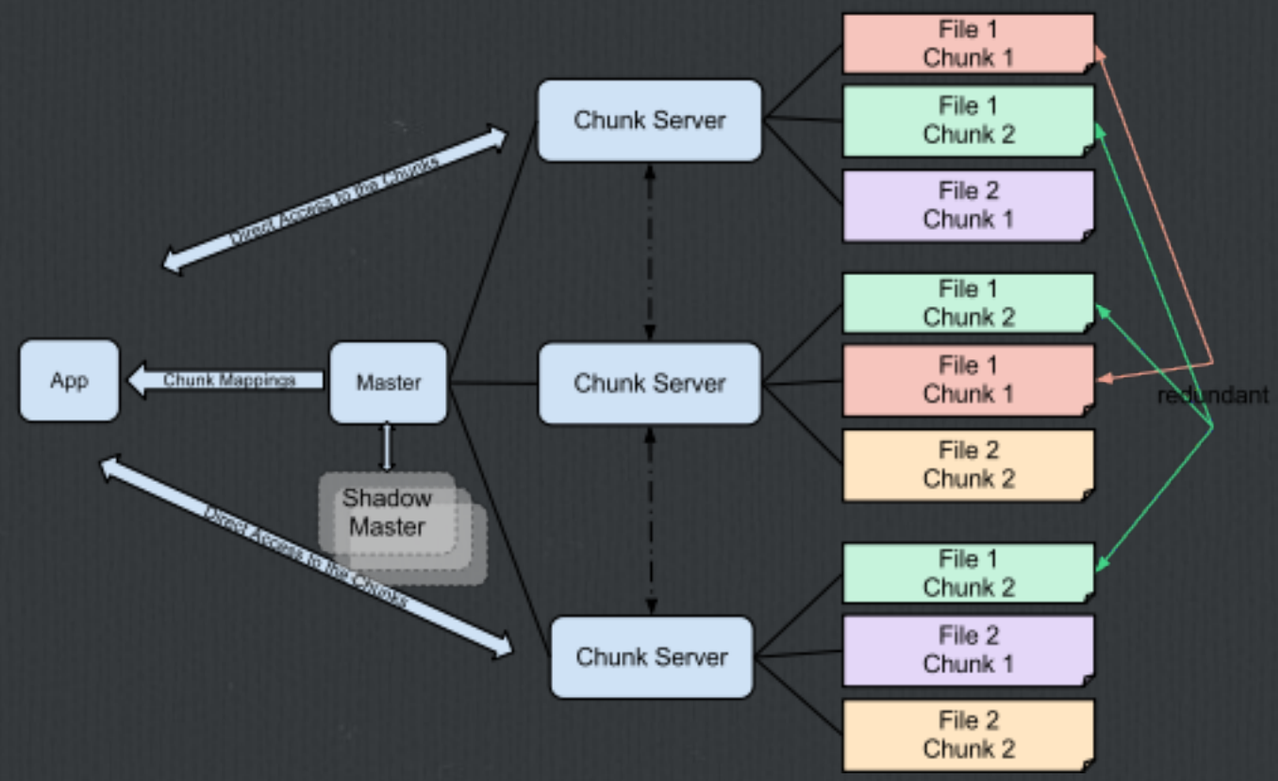

Google File System

MapReduce

Two Separate Tasks:

Map - takes a set of data and breaks it down into tuples (key/value pairs) to be distributed to worker nodes.

Reduce - takes the output from the workers and aggregate the result.

MapReduce: Analogy

Roman Times Census:

- Map: Census bureau would dispatch its people to each city in the empire. Each census taker in each city would be tasked to count the number of people in that city and then return their results to the capital city.

- Reduce: Aggregate all results to a single count (sum of all cities) to determine the overall population of the empire.

Motivation: For Hadoop

- ☐ Google's MapReduce and Filesystem are proprietary
- ☐ Hadoop is opensource software by Apache (you can use the software for free!)
- ☐ Not only is it opensource, the community is great!

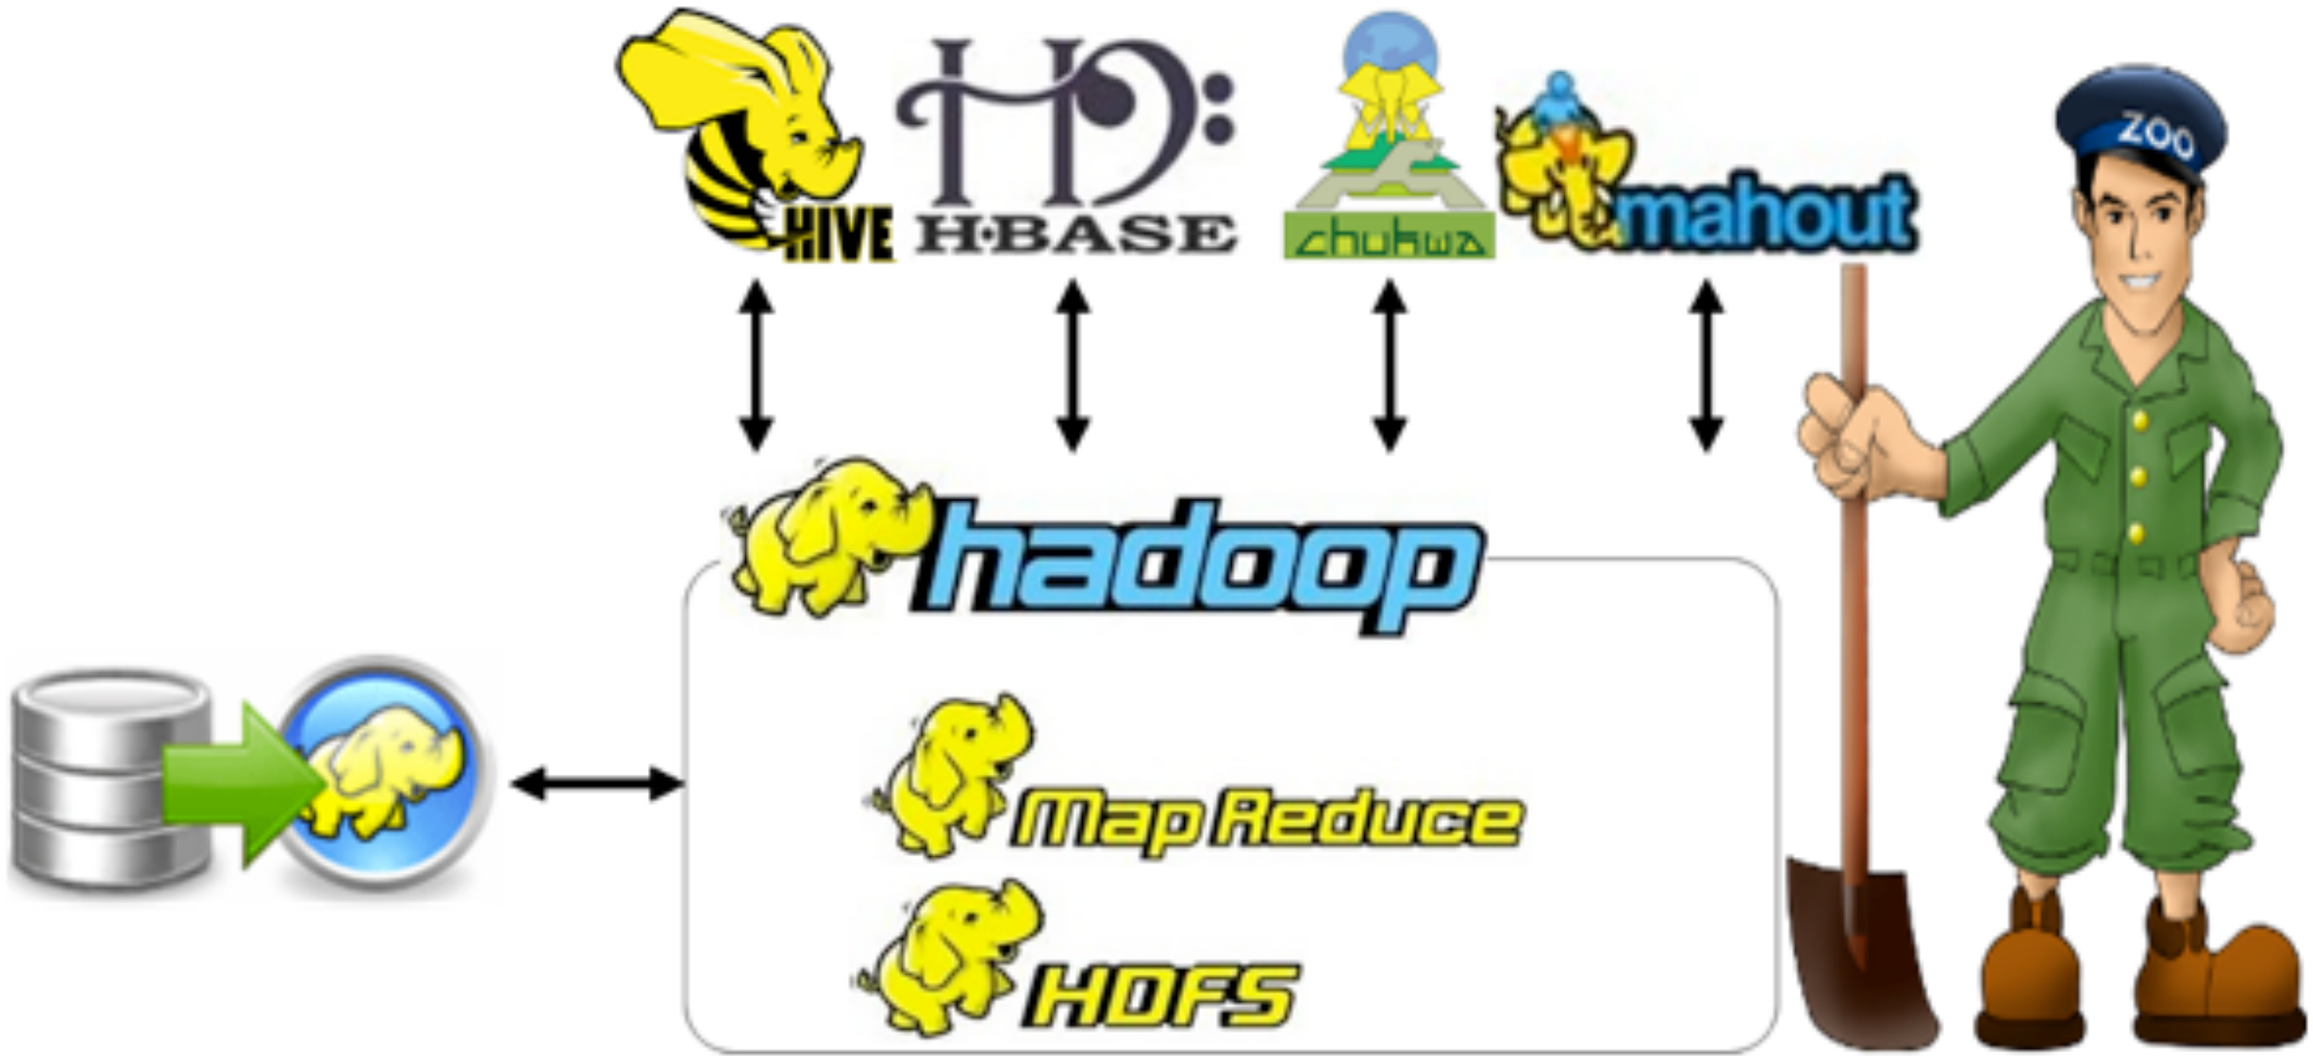

Hadoop Ecosystem

Hadoop, Hive, HDFS, MapReduce, Mahout, HBase, Zookeeper, Chukwa, Pig

HBase

- distributed database modeled after Google's BigTable (which is built on top of the Google Filesystem)
- written in Java (but has wrappers for other languages)
- fault-tolerant
- good for sparse data

Hive

- SQL-like language for accessing HDFS

Pig

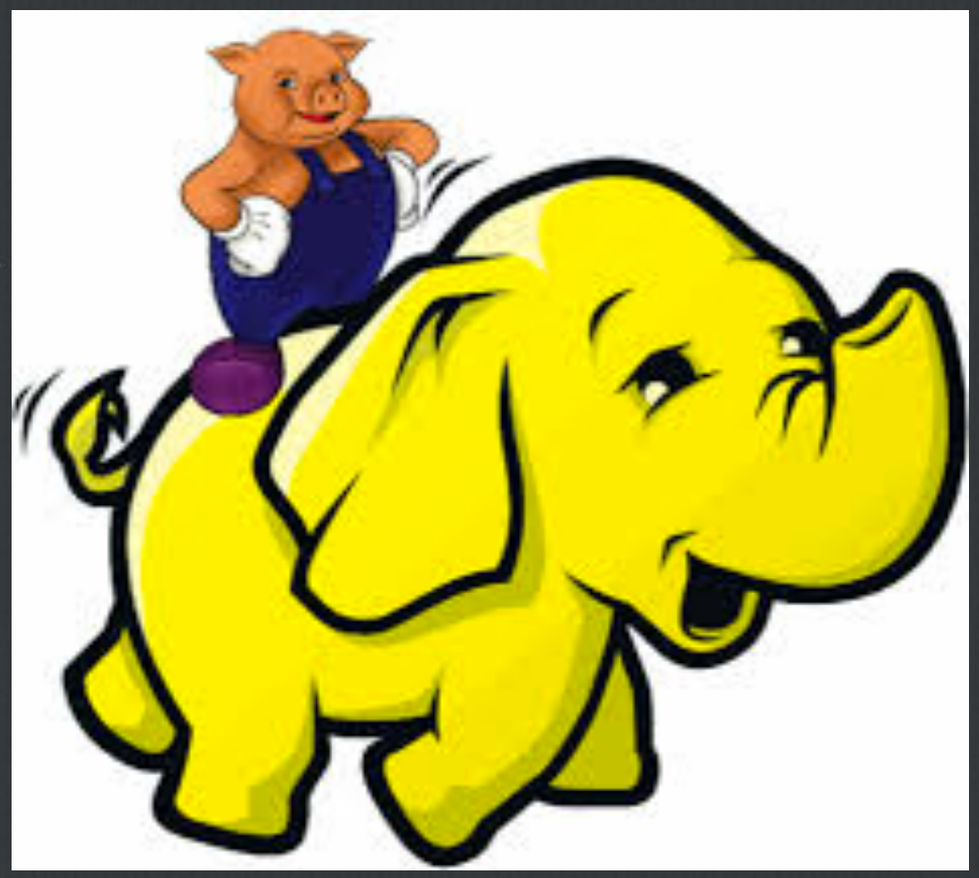

For data extraction like Hive, but has its own language: Pig Latin

Unlike Hive it can:

- ☐ use lazy evaluation (delays the evaluation of an expression until its value is needed)
- ☐ use ETL (Extract, Transform, Load)
- ☐ store data at any point during a pipeline
- ☐ declare execution plan
- ☐ support pipeline splits

Other

- [ ] Chukwa - data collection system for monitoring large distributed systems
- [ ] Mahout - scalable machine learning libraries
- [ ] Zookeeper - service for maintaining clusters

Ways you may already use Hadoop

- ☐ HBase powers Facebook Messenger
- ☐ Yahoo! search is processed by Hadoop
- ☐ Ebay uses Hadoop for search optimization and research
- ☐ Hulu uses Hadoop and HBase for storage

And a ton of other products…

WHY: Hadoop on AWS

- ☐ Elastic
- ☐ Cheap
- ☐ Easier than doing DevOps yourself

AWS Basics

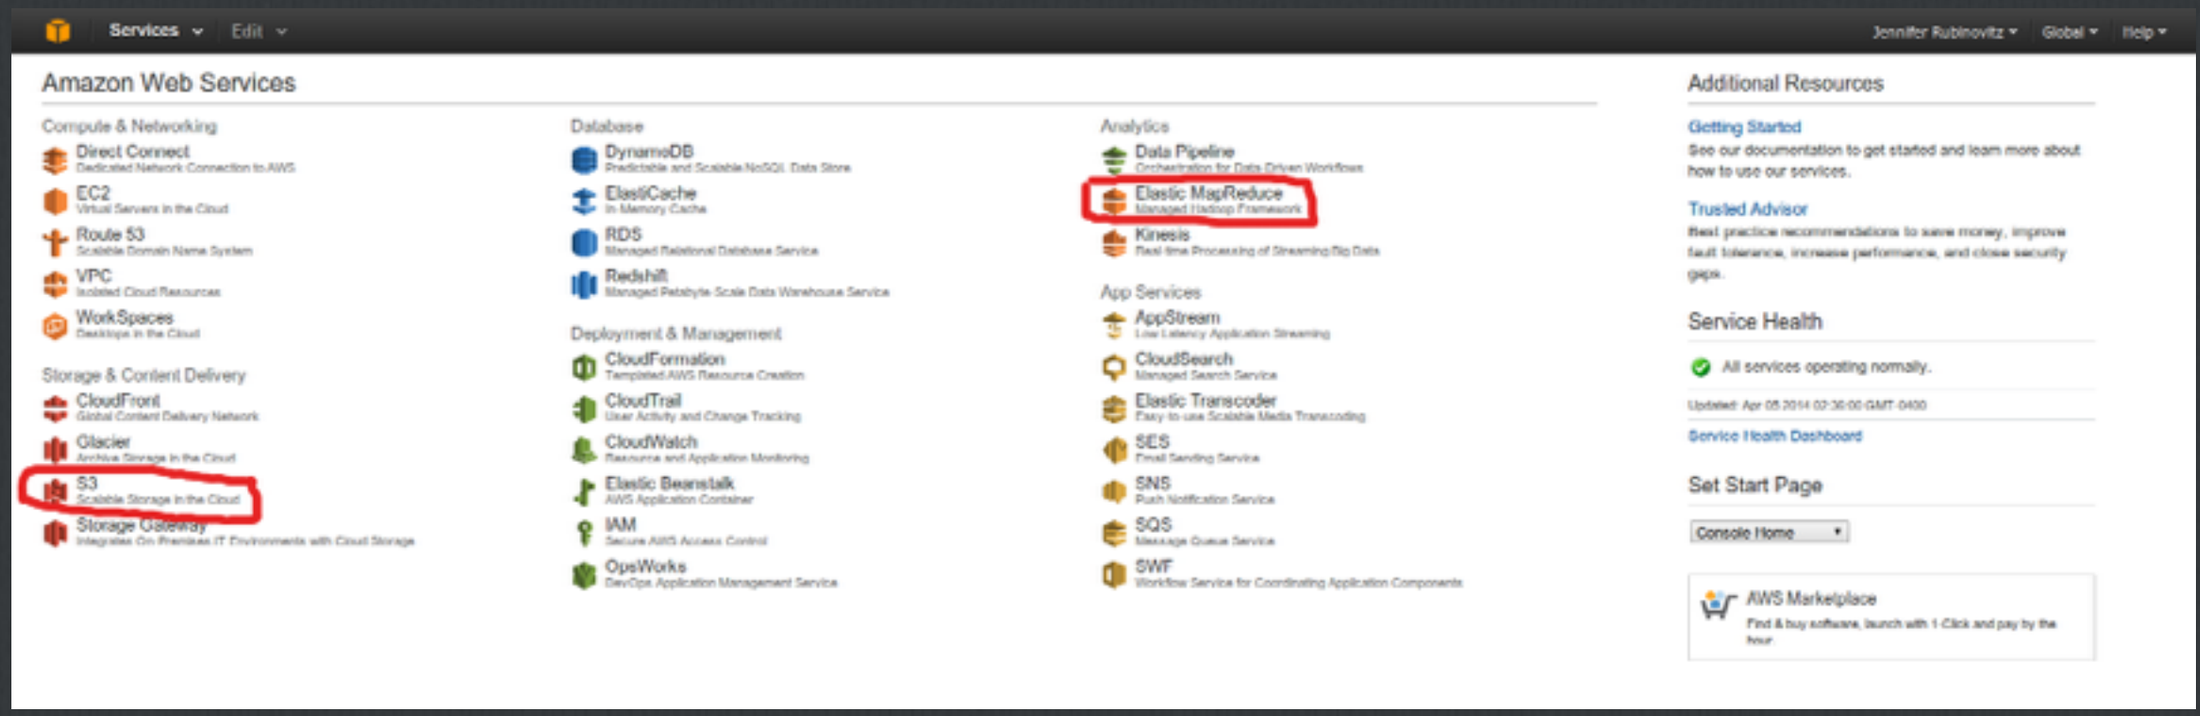

- [ ] Go to console.aws.amazon.com

HOW: Hadoop on AWS

- [ ] AWS Console GUI (www.console.aws.amazon.com)
- [ ] Ruby MapReduce Command Line Interface (http://aws.amazon.com/developertools/Elastic-MapReduce/2264): requires Ruby 1.8.7!!

Our Example

The Hello World of Hadoop is…

Word counting

Steps:

1. Write a mapper in Python

2. Put it into AWS S3

3. Launch an EMR instance

wordSplitter.py

```
#!/usr/bin/python

import sys

import re


def main(argv):

    pattern = re.compile("[a-zA-Z][a-zA-Z0-9]*")

   for line in sys.stdin:

        for word in pattern.findall(line):

            print "LongValueSum:" + word.lower() + "\t" + "1"


if __name__ == "__main__":

    main(sys.argv)
```

Step 2: Setup a S3 Bucket For Storage

- ☐ Go to https://console.aws.amazon.com/s3/
- ☐ Click “Create Bucket”
- ☐ Create a bucket to keep your data

Step 3: Setup Cluster

- [ ] Go to https://console.aws.amazon.com/elasticmapreduce/
- [ ] Click “Create Cluster”
- [ ] or in the CLI
- [ ] ./elastic-mapreduce —create —stream —mapper s3://<our-bucket>/wordSplitter.py —output s3://<our-bucket>/output —reducer aggregate

Step 4: Wait

- [ ] Go running
- [ ] Paint your nails
- [ ] Read a book
- [ ] Do other work? Naaa.

Example: Google Ngrams

- AWS has the 2 TB dataset of Ngrams in books over time
- We can use Hive to query them and find trends
- Dataset used for http://books.google.com/ngrams

Setup a hive instance

$ ./elastic-mapreduce —create —alive —hive-interactive

$ ./elastic-mapreduce —list <job-flow-id>

$ ./elastic-mapreduce —ssh <job-flow-id>

Setup Hive Tables

```
$ hive

$ set hive.base.inputformat=org.apache.hadoop.hive.ql.io.HiveInputFormat;

$ set mapred.min.split.size=134217728;

$ CREATE EXTERNAL TABLE english_1grams (
 gram string,
 year int,
 occurrences bigint,
 pages bigint,
 books bigint
)
ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t'
STORED AS SEQUENCEFILE
LOCATION 's3://datasets.elasticmapreduce/ngrams/books/20090715/eng-all/1gram/';
```

Normalize the Data

(convert to lowercase and ignore extraneous characters)

```
CREATE TABLE normalized (
 gram string,
 year int,
 occurrences bigint
);
INSERT OVERWRITE TABLE normalized
SELECT
 lower(gram),
 year,
 occurrences
FROM
 english_1grams
WHERE
 year >= 1890 AND
 gram REGEXP "^[A-Za-z+'-]+$";
```

Word-ratio by Decade

```
CREATE TABLE by_decade (

 gram string,

 decade int,

 ratio double

);
```

```
INSERT OVERWRITE TABLE by_decade
SELECT
 a.gram,
 b.decade,
 sum(a.occurrences) / b.total
FROM
 normalized a
JOIN (
 SELECT
  substr(year, 0, 3) as decade,
  sum(occurrences) as total
 FROM
  normalized
 GROUP BY
  substr(year, 0, 3)
) b
ON
 substr(a.year, 0, 3) = b.decade
GROUP BY
 a.gram,
 b.decade,
 b.total;
```

Results

1900

radium, ionization, automobiles, petrol, archivo, automobile, electrons, mukden, anopheles, marconi, botha, ladysmith, lhasa, boxers, suprema, aboord, rotor, turkes, wireless, conveyor, manchurian, erythrocytes, shoare, thirtie, kop, tuskegee, thorium, audiencia, bvo, arteriosclerosis

1910

cowperwood, britling, boches, montessori, venizelos, bolsheviki, salvarsan, photoplay, pacifists, joffre, petrograd, pacifist, bolshevism, airmen, kerensky, foch, boche, serbia, serbian, hindenburg, madero, serbians, bombing, ameen, anaphylaxis, aviators, syndicalism, aviator, biplane, taxi

…

1930

dollfuss, goebbels, manchukuo, hitler, sudeten, hitler's, rearmament, nazis, wpa, nazi, nra, manchoukuo totalitarian, pwa, tva, stalin's, peiping, homeroom, kulaks, stalin, devaluation, bta, carotene, broadcasts, corporative, comintern, ergosterol, reichswehr, ussr, businessmen

…

2000

bibliobazaar, itunes, cengage, qaeda, wsdl, aspx, xslt, actionscript, xpath, sharepoint, blogs, easyread, ipod, xhtml, blog, rfid, google, writeline, proteomics, bluetooth, voip, microarray, mysql, microarrays, putin, dreamweaver, dvds, ejb, xml, osama

Next Steps

- [ ] Go to www.github.com/rubinovitz/hackny-masters-hadoop for more examples and reference
- [ ] Questions?